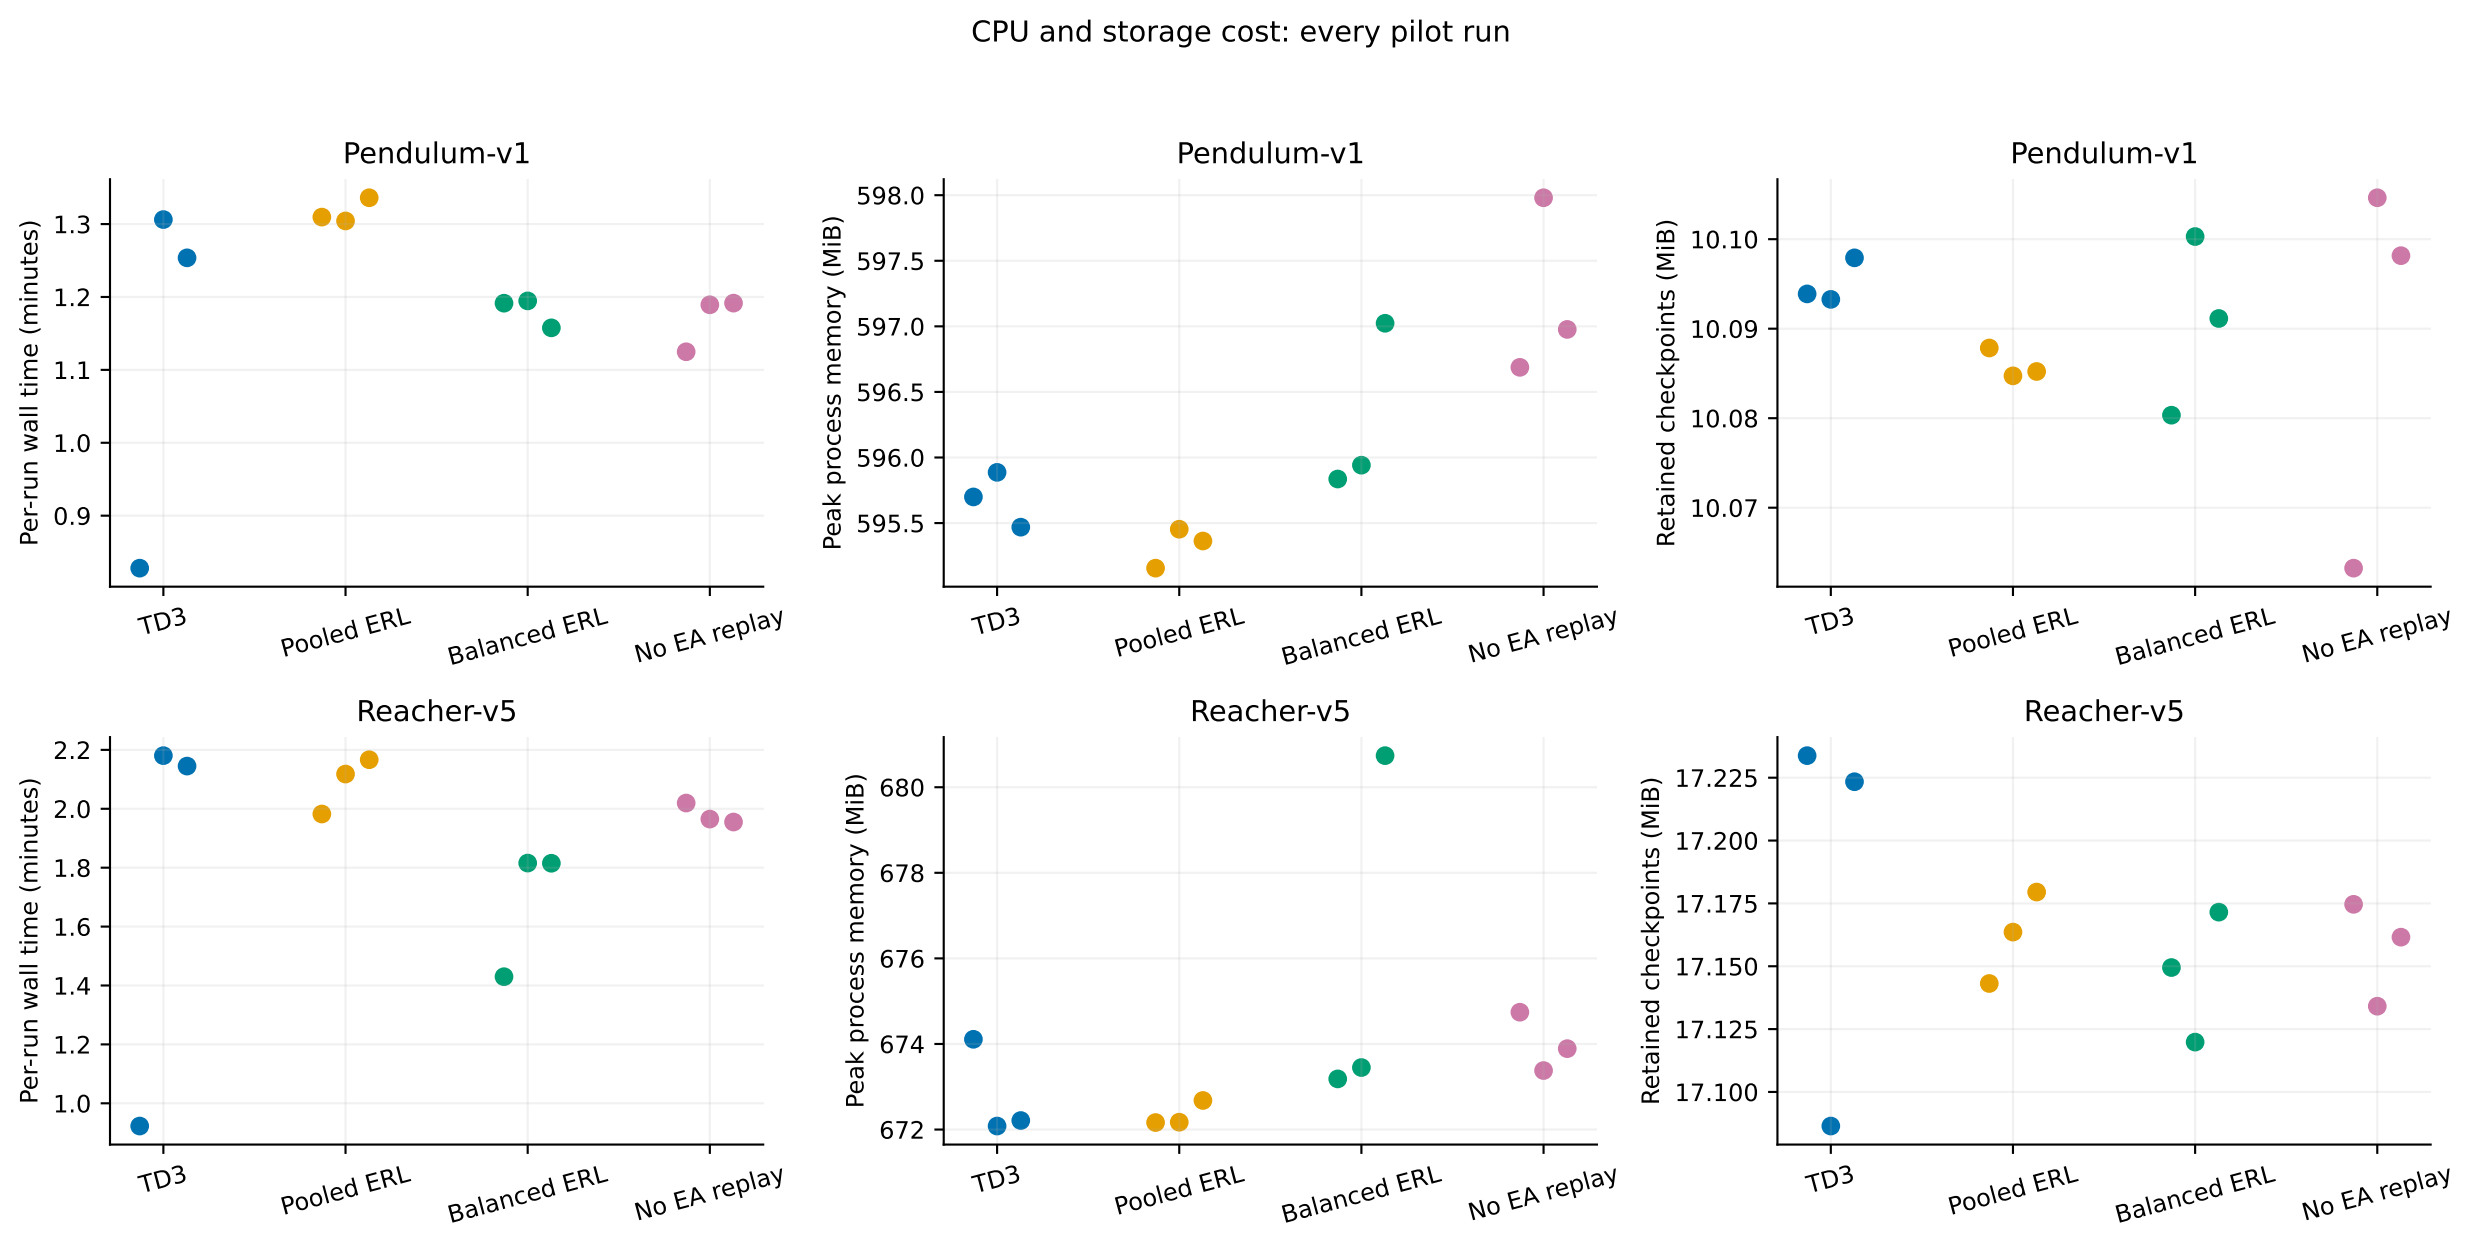

Each point is one run, including evaluation and checkpoint writes. The two seed-zero TD3 runs were profiled with 1 worker(s), then reused in the 24-run pilot; the remaining runs used 2 concurrent worker(s). Each run used one computational thread. Different contention makes per-method wall-time rankings unfair. Wall times overlap; peak RSS is per process.